

LawBot

An Engineering Project in Community Service

Phase – II Report

Submitted by

Anushka Dubey (23BCE11492)

Dev Anurag Varshney (23BCE10880)

Aarul Kumar (23BCE10857)

Gauri Makker (23BCE11131)

Vanshika Dhaka (23BCE11463)

Udita Chakraborty (23BAI10653)

in partial fulfillment of the requirements for the degree of

Bachelor of Technology



VIT[®]
BHOPAL
www.vitbhopal.ac.in

VIT Bhopal University
Kothrikalan, Sehore
Madhya Pradesh

March, 2026



Bonafide Certificate

Certified that this project report titled “LawBot” is the Bonafide work of “23BCE10857 Aarul Kumar, 23BCE11492 Anushka Dubey, 23BCE10880 Dev Anurag Varshney, 23BCE11131 Gauri Makker, 23BCE11463 Vanshika Dhaka, 23BAI10653 Udit Chakraborty” who carried out the project work under my supervision.

This project report (Phase II) is submitted for the Project Viva-Voce examination held on March 30, 2026.

Dr Lakshmi D
(Supervisor)

Dr. Manisha Jain
(Reviewer 1)

Dr. Priscilla Dinkar Moyya
(Reviewer 2)



Declaration of Originality

We, hereby declare that this report entitled **LawBot** represents our original work carried out for the EPICS project as a student of VIT Bhopal University and, to the best of our knowledge, it contains no material previously published or written by another person, nor any material presented for the award of any other degree or diploma of VIT Bhopal University or any other institution. Works of other authors cited in this report have been duly acknowledged under the section "References".

Date

March 30, 2026

Reg No & Name

23BCE11492 Anushka Dubey

23BCE10880 Dev Anurag Varshney

23BCE10857 Aarul Kumar

23BCE11131 Gauri Makker

23BCE11463 Vanshika Dhaka

23BAI10653 Udit Chakraborty

Acknowledgement

This work, has benefited in various ways from several people. Whilst it would be simple to name them all, it would not be easy to thank them enough. First and foremost, we express our profound gratitude to our project supervisor, Prof. Dr Lakshmi D, for her invaluable guidance, unwavering support, and insightful feedback throughout the development of this project. Her mentorship was instrumental in shaping our ideas and steering this project in the right direction. We are also deeply thankful to our project reviewers, Dr. Manisha Jain and Dr. Priscilla Dinkar Moyya, for their time, critical evaluation, and constructive suggestions during the review phases. Their comments and observations have been crucial in enhancing the quality and depth of our work. We extend our sincere thanks to VIT Bhopal University for providing us with the EPICS platform and the necessary resources to undertake this project. Our appreciation also goes to our peers and the student community who participated in our pilot testing, providing us with essential user feedback. Lastly, we would like to acknowledge the support and encouragement of our families and friends, whose constant belief in us gave us the strength to persevere and complete this project.

Abstract

Access to legal information and effective grievance redressal mechanisms remains a persistent challenge for a significant proportion of the population in rural India. This issue is exacerbated by a shortage of qualified legal professionals, financial limitations, low levels of legal literacy, and the inherent complexity of administrative and judicial procedures. Consequently, many individuals are unable to fully understand their legal rights, initiate formal complaints, or obtain timely and appropriate legal assistance.

This paper presents LawBot, an AI-driven legal assistance system designed to enhance access to justice through scalable and inclusive technological interventions. The proposed system integrates advanced methodologies, including Large Language Models (Llama 3), Retrieval-Augmented Generation (RAG), and vector-based semantic search using Pinecone, to deliver contextually relevant and legally grounded responses derived from verified Indian legal sources. By combining generative capabilities with retrieval-based validation, the system aims to mitigate issues related to hallucination and ensure informational reliability.

To address barriers related to literacy and digital accessibility, LawBot incorporates Whisper-based speech recognition, enabling voice-driven interaction in multiple languages. Furthermore, deployment via widely adopted communication platforms such as WhatsApp facilitates broad reach and usability, particularly in resource-constrained and rural settings. The system provides conversational legal guidance, translates complex legal terminology into simplified language, assists users in drafting basic grievance and complaint documents, and offers procedural support for navigating formal redressal mechanisms.

By adopting a user-centric and technology-enabled approach, LawBot contributes to bridging critical gaps in legal accessibility. The system aligns with broader national objectives of digital inclusion and citizen empowerment, demonstrating the potential of AI-driven solutions to reduce socio-economic disparities, improve legal awareness, and strengthen participatory access to justice.

Index

Sl. No.	Topic	Page No.
1.	Introduction	10
1.1	Motivation	10
1.2	Problem Statement	11
1.3	Objectives	12
1.4	Scope of the Study	13
2.	Literature Review	15
2.1	Traditional Legal Aid Systems	15
2.2	Online Grievance Redressal Portals	15
2.3	Legal Literacy Initiatives	16
2.4	AI-based Legal Assistance Systems (Global Context)	17
2.5	RAG-Based Legal Assistants in India	19
2.6	Limitations of Existing Systems	20
2.7	Research Gap Identified	21
3.	Proposed System: LawBot	24
3.1	System Overview	24
3.2	System Architecture	24
3.3	Technology Stack	27
3.3.1	Large Language Model (Llama 3)	27
3.3.2	Retrieval-Augmented Generation (RAG)	27
3.3.3	Vector Database (Pinecone)	29
3.3.4	Speech Processing (Whisper)	30
3.3.5	Deployment Platform (WhatsApp Integration)	30
3.4	Data Collection and Preprocessing	31
3.5	Working Principle	33
3.6	System Workflow	34
3.7	Use Case Scenarios	36
4.	Implementation and Results	42
4.1	System Implementation	42
4.2	Experimental Setup	42
4.3	Performance Evaluation	43
4.4	Results and Analysis	44

4.5	Limitations of the System	46
5.	Discussion	47
5.1	Comparison with Existing Systems	47
5.2	Practical Implications	50
6.	Conclusion and Future Work	52
6.1	Conclusion	52
6.2	Future Enhancements	52
7.	References	55
8.	Appendix	57
8.1	Sample Outputs / Screenshots	57
8.2	Questionnaire / User Feedback	60
9.	Biodata	65

List of Figures

Sl. No	Caption	Page No
1	System architecture of LawBot	26
2	Sequence diagram of LawBot interaction (simple query).	35
3	Experiment pipeline (sequence diagram).	45
4	Integration points of LawBot with existing systems.	48
5	LawBot home screen	58
6	LawBot chatbot (RAG + LLM + AI Legal Aid)	59
7	LawBot document ingestion	59
8	LawBot response	60
9	User Satisfaction Distribution	62
10	Ease of Use Distribution	63
11	User Recommendation	63
12	Response Clarity	64

List of Table

Sl. No	Caption	Page No
1	Comparison of Key Legal Assistance Systems	22
2	Implementation Options Comparison	39
3	Evaluation metrics (mean $\pm$ 95% CI). Retrieval precision and recall, BLEU score, and latency are shown.	43
4	Comparison of existing system	48
5	LawBot survey data	61

INTRODUCTION

Access to legal support is a fundamental right; however, a significant proportion of citizens in India—particularly those residing in rural and underserved regions—face substantial challenges in understanding legal frameworks, seeking guidance, and accessing grievance redressal mechanisms. Factors such as limited availability of legal professionals, financial constraints, geographical barriers, and low levels of legal literacy contribute to a widening gap between citizens and the justice delivery system (National Legal Services Authority, 2022).

In recent years, initiatives under the Digital India program have improved digital connectivity and access to online services (Ministry of Electronics and Information Technology, 2023). Simultaneously, advancements in Artificial Intelligence (AI), especially in Natural Language Processing (NLP), have enabled the development of intelligent systems capable of understanding and generating human-like language (Brown et al., 2020). These developments present an opportunity to design scalable, user-friendly solutions that can enhance legal accessibility.

This project introduces *LawBot*, an AI-powered legal assistance system designed to democratize access to legal information and grievance support. LawBot provides simplified, accurate, and multilingual legal guidance through a conversational interface, making it accessible to users with varying levels of literacy. The system leverages advanced technologies, including Large Language Models (Touvron et al., 2023), Retrieval-Augmented Generation (Lewis et al., 2020), vector-based semantic search using Pinecone (Pinecone Systems Inc., 2023), and Whisper-based speech recognition (Radford et al., 2023).

By integrating these technologies, LawBot ensures that responses are contextually relevant and grounded in verified legal sources. Furthermore, deployment through widely used platforms such as WhatsApp enhances accessibility, particularly in rural areas where mobile penetration is high (Telecom Regulatory Authority of India, 2023). The system enables users to ask legal questions, understand complex legal terminology, draft basic grievance documents, and receive structured guidance for complaint filing.

1.1 Motivation

The development of LawBot is motivated by persistent socio-economic and systemic challenges that limit access to legal services in India. Despite being one of the world’s largest democracies, India faces a significant disparity in the distribution of legal resources, particularly between urban and rural areas (World Bank, 2022).

A major concern is the shortage of legal professionals in rural regions, which restricts timely access to legal advice. Financial barriers further exacerbate the problem, as engaging with the legal system involves expenses such as travel, documentation, and consultation fees, which are often unaffordable for low-income populations (NITI Aayog, 2021).

Another critical issue is low legal literacy. A large segment of the population lacks awareness of fundamental rights, government schemes, and grievance redressal platforms such as CPGRAMS and RTI (Government of India, 2022). Even when awareness exists, individuals often struggle to understand legal terminology and procedural requirements.

Technological limitations also persist. While several online grievance portals exist, they are not always user-friendly for individuals with limited digital literacy. In contrast, mobile-based platforms, particularly messaging applications like WhatsApp, have seen widespread adoption across rural India (TRAI, 2023), making them an ideal medium for delivering legal assistance.

Advancements in AI further strengthen the motivation for this project. Large Language Models (LLMs) have demonstrated strong capabilities in natural language understanding and generation (Brown et al., 2020). When combined with Retrieval-Augmented Generation (Lewis et al., 2020), these models can provide reliable and factually grounded responses. Additionally, speech recognition systems such as Whisper enable voice-based interaction, improving accessibility for low-literacy users (Radford et al., 2023).

1.2 Problem Statement

Despite the presence of well-established legal frameworks and digital grievance redressal systems in India, a significant portion of the population—particularly in rural and underserved regions—remains unable to effectively access these services. This gap arises due to a combination of socio-economic, linguistic, and technological barriers.

One of the primary challenges is the limited availability of legal professionals in rural areas, which restricts timely access to expert guidance (National Legal Services Authority, 2022). Additionally, the high cost of legal consultation and court-related processes discourages economically weaker sections from seeking formal legal support (NITI Aayog, 2021).

Another critical issue is the complexity of legal language and procedures, which makes it difficult for non-experts to understand their rights or navigate grievance mechanisms. Many citizens also lack awareness of available platforms such as CPGRAMS and RTI, further limiting their ability to seek justice (Government of India, 2022).

Although digital platforms have been introduced to improve accessibility, they often require a certain level of digital literacy and are not designed for conversational or intuitive interaction. This creates an additional barrier for users with limited education or technological exposure.

Therefore, the problem addressed in this project can be formulated as:

“To design and develop an intelligent, scalable, and user-friendly system that provides accurate legal information, simplifies complex legal procedures, and enables citizens—especially in rural areas—to access legal assistance efficiently.”

The system must ensure:

- Accessibility for users with low literacy and limited resources
- Accuracy of legal information through verified sources
- Usability via a simple conversational interface
- Scalability to serve a large and diverse population

1.3 Objectives

The primary objective of the LawBot system is to develop an AI-powered legal assistance platform that enhances accessibility, understanding, and usability of legal information for a wide range of users. The project aims to leverage modern AI technologies to address existing gaps in the legal support ecosystem.

The specific objectives are outlined as follows:

- **Provide Accessible Legal Guidance:**
To deliver instant, accurate, and easy-to-understand responses to user queries related to legal rights, government schemes, and grievance procedures.
- **Simplify Complex Legal Language:**
To translate technical legal terminology into simple and comprehensible language while preserving the original meaning and intent.
- **Assist in Grievance Filing and Document Preparation:**
To support users in drafting structured complaint letters, applications, and RTI requests based on their specific issues.
- **Enable Multilingual and Voice-Based Interaction:**
To allow users to interact with the system in multiple languages and through voice input using speech recognition technologies (Radford et al., 2023).
- **Ensure Response Accuracy through RAG:**

- To improve reliability by grounding responses in verified legal documents using Retrieval-Augmented Generation techniques (Lewis et al., 2020).
- Leverage Widely Used Communication Platforms:
To deploy the system on platforms such as WhatsApp, ensuring high accessibility and ease of adoption (Telecom Regulatory Authority of India, 2023).
- Reduce Dependency on Legal Experts for Basic Queries:
To empower users with sufficient knowledge to handle basic legal issues independently, reducing unnecessary reliance on legal professionals.
- Promote Digital Inclusion and Legal Awareness:
To align with national initiatives such as Digital India by making legal services more accessible through digital means (Ministry of Electronics and Information Technology, 2023).
- Ensure Scalability and Adaptability:
To design a flexible system capable of expanding across different legal domains, regions, and languages with minimal modifications.

1.4 Scope of the Study

The scope of this project focuses on the design and implementation of an AI-based legal assistance system that provides general legal guidance and supports users in navigating grievance redressal processes.

The system is intended to:

- Provide basic legal information related to rights, procedures, and government schemes
- Assist users in filing common grievances, such as consumer complaints and public service issues
- Enable drafting of simple legal documents, including complaint letters and RTI applications
- Support multilingual communication (e.g., Hindi and English)
- Offer voice-based interaction to accommodate users with low literacy levels
- Deliver services through accessible platforms such as WhatsApp, ensuring ease of use

However, the system operates within certain limitations:

- It does not replace professional legal advisors or lawyers
- It does not handle complex legal cases or litigation-specific guidance
- It provides informational support only and not legally binding advice

Despite these constraints, LawBot serves as an effective tool for improving legal awareness, accessibility, and initial-level support. The system is particularly beneficial for underserved populations, helping bridge the gap between citizens and the legal ecosystem while promoting informed participation in grievance redressal processes.

Literature Review

2.1 Traditional Legal Aid Systems

Traditional legal aid in India is structured under the Legal Services Authorities Act (1987), which established NALSA (National Legal Services Authority) and state/district bodies to provide free legal aid to the disadvantaged. This network includes Lok Adalats for mediation and cadre of Para-legal Volunteers (PLVs) who conduct outreach. Despite its broad framework, the system faces severe resource constraints. As of 2023, there are 49,126 panel lawyers for legal aid, a ratio of only one lawyer for ~12 villages. PLVs number 42,446 (~1 per 14 villages), indicating that the majority of villages have no immediate legal contact. The rural population (~68% of India) is disproportionately underserved. These personnel shortfalls manifest in long travel distances (many litigants travel 50–300 km for hearings) and delayed redress.

Examples: Gram Sabhas and Nari Adalats (women's courts) are informal dispute-resolution forums recognized by law (e.g. under Bihar's Panchayat Raj Act). These grassroots efforts aid access but address only certain disputes (e.g. matrimonial, land). The tele-legal initiative (Tele-Law) uses CSC video-conferencing to connect villagers with lawyers, yet requires local CSC infrastructure and basic literacy.

Strengths: Formal legal aid guarantees counsel for eligible individuals (e.g. Scheduled Castes/Tribes, women, minors) under Article 39A. Lok Adalats resolve cases quickly with zero court fees. Widespread outreach (e.g. legal literacy camps, Lok Adalat circuits) has increased awareness: in 2022 over 41,000 literacy programs were conducted nationwide.

Limitations: Coverage is uneven—lawsuit pressures are far higher in cities than in villages. The ratio of mobile legal aid vehicles is only 1 per 7,600+ villages, so rural citizens have minimal physical access. Legal processes remain complex in language and procedure; 70% of rural residents are illiterate in legal matters. Consequently, many rural citizens rely on informal/unverified advice.

Implications for LawBot: The shortage of in-person aid and low legal literacy underscore the need for an automated, scalable system. LawBot must operate through ubiquitous means (mobile chat) and simplify legal jargon, to supplement human services and extend reach to isolated communities.

2.2 Online Grievance Redressal Portals

The Government of India has established several online platforms to lodge and track complaints. The central portal CPGRAMS (Centralized Public Grievance Redress and Monitoring System) allows citizens to file grievances 24×7 via web or mobile app. Users receive a unique registration ID to track status and can appeal unsatisfactory responses. State governments have analogous portals (e.g. Rajasthan's Lok Taknik etc.), and sector-specific systems exist (RTI online, consumer forums, etc.).

Examples: The Controller General of Accounts notes that CPGRAMS connects all Union and State Ministries/Departments, with mobile apps (Google Play, UMANG) to broaden reach. These portals triage complaints by category and enforce response timelines.

Technical Mechanisms: Portals use centralized databases and workflow engines. They are purely text-based, sometimes supporting attachments (photos, documents). They provide transparency but no conversational guidance.

Strengths: They eliminate intermediaries and bureaucratic layers, enabling citizens to bypass corrupt or slow officials. Data analytics on grievance trends informs policy-making. The accessibility of platforms (web + app) means no physical travel is needed, benefiting remote citizens.

Limitations: Online portals assume internet literacy. Many rural users find the interfaces complex (forms, dropdowns). There is no interactive help: users must know which department to address, and how to articulate their issue. Voice or local-language support is generally absent. Also, digital verification (OTP, email) can be barriers in areas with poor connectivity or shared phones.

Implications for LawBot: While existing portals solve the submission problem, LawBot's RAG-based agent would enhance them by offering conversational guidance on which portal to use, how to word complaints, and what legal basis applies. Integrating with WhatsApp (already widespread) circumvents much of the digital literacy barrier.

2.3 Legal Literacy Initiatives

Recognizing that awareness is key to access, the legal services apparatus runs extensive literacy and awareness programs. For example, NALSA and State LSAs conduct seminars, street plays (nukkad nataks), pamphleteering, radio/TV programs, and mobile legal van campaigns to explain rights. Since the mid-2010s, India has also focused on education: ~18,745 Legal Literacy Clubs operate in schools nationwide, and in 2022 over 41,224 awareness programs were held.

Annual funding for these activities has grown (Rs.145 crore in 2021–22; Rs.190 crore in 2022–23).

Key Examples:

- Legal Literacy Clubs (Rajasthan & national): Court-run clubs in schools encourage students to discuss legal topics. This has scaled to nearly 20,000 schools.
- Pan-India Campaigns: Periodic national drives (e.g. “Nyaya Jagriti”) mobilize volunteers to visit villages with mobile vans and group lectures.
- Lok Adalats and Legal Aid Clinics: These are forums for both dispute resolution and community outreach, often coupled with awareness workshops on-site.

Strengths: These initiatives increase general awareness of basic rights (e.g. right to information, schemes like BPL benefits) and legal remedies (e.g. lodging FIRs, filing RTI requests). They also publicize free legal aid availability, improving trust in the system. Post-implementation surveys show two-thirds of rural respondents know about free legal aid services.

Weaknesses: Such campaigns are intermittent and one-directional. Awareness events do not provide real-time query response or personalized assistance. They generally use Hindi or English, limiting impact on speakers of tribal/local languages. Reaching remote areas remains costly: e.g. only 78 mobile legal vans exist nationwide, one per ~7,600 villages. There is no scalable on-demand mechanism.

Implications for LawBot: LawBot can be viewed as a digital complement to these efforts. Its conversational AI can replicate some outreach functions continuously (answering queries any time) and provide language-appropriate simplifications. Data from literacy programs (e.g. frequent questions) can inform LawBot’s knowledge base, while LawBot can circulate digital pamphlets or voice messages to further legal education.

2.4 AI-based Legal Assistance Systems (Global Context)

Worldwide, AI and chatbots have begun to transform access to legal information. High-profile examples include DoNotPay (USA), Ross Intelligence (USA), Modria (USA), and Ailira (Australia). More recently, general-purpose LLMs (ChatGPT, Google Bard, Bing Chat) have been repurposed as ad hoc legal advisors.

Key Examples:

- DoNotPay (founded 2015) bills itself as “the world’s first robot lawyer,” initially helping users contest parking tickets and cancel subscriptions. It uses rule-based logic combined with NLP to automate routine tasks. Notably, it sparked legal

debates: courts in New York ruled it was not practicing law, but others sued it for unauthorized legal advice. Its business model is subscription-based, providing document templates and chatbot Q&A for consumer and small claims law.

- ROSS Intelligence (founded 2015) created an AI search engine for legal research using IBM Watson. It let lawyers query case law in natural language. ROSS shut down in 2020 after patent disputes, but it demonstrated LLM potential in legal research.
- ChatGPT (OpenAI) and Google Bard: Though not legal-specific, these LLMs are globally accessible. They have been widely tested for legal queries. Users find them cost-effective and fast (seconds to answer). However, studies show frequent hallucination (fabricated “answers”) and incomplete coverage. For example, legal experts observed that GPT-3.5/4 produced legally incorrect advice in ~20% of cases, raising malpractice concerns.
- Tessa (UK) – launched by a UK legal aid charity (LawWorks), Tessa uses an LLM fine-tuned on UK law to assist vulnerable clients. It provides triage and signposting in simple language (citation [Stockton 2023], see below).

Technical Mechanisms: These systems typically combine NLP understanding of user questions with knowledge bases (statutes, cases, FAQs). Pure LLMs generate responses from internal probabilities, often without guaranteeing factual grounding. Retrieval-augmented systems (RAG) link an LLM to a database: the user’s query is matched to relevant documents (via embeddings or keyword search), and the LLM conditions its answer on the retrieved text. Pinecone or FAISS are common vector databases used in RAG.

Strengths: AI legal assistants can dramatically lower cost barriers. Stanford’s Justice Innovation Lab notes that chatbots “enable cost-effective access to information” and can reach non-English speakers if multilingual models are used. AI can operate 24/7, handle large query volumes, and adapt rapidly to new laws via updated data.

Weaknesses: According to multiple sources, major challenges include:

- Accuracy & Hallucination: LLMs occasionally invent statutes or case references. This is especially risky in legal contexts. The Stanford report warns of misinformation and the danger of unauthorized practice of law.
- Bias & Fairness: Models trained on internet data may encode biases (e.g. gender/race bias in legal scenarios), raising ethical issues (DoNotPay’s IPO scandal is one example of bias/algo malpractice).
- Regulation & Trust: Legal advice is a regulated profession. AI systems must be carefully framed as tools, not licensed attorneys. The Stanford project emphasizes the need for clear disclaimers and user education.

- **Language/Locality:** Many AI solutions focus on English and Western law. Non-English speakers or Indian law users find fewer relevant systems.

Implications for LawBot: Global lessons suggest LawBot must prioritize grounding answers in official Indian legal texts (to avoid hallucination) and operate in Hindi/vernacular. Voice interaction is especially critical: users may prefer speaking to a bot rather than typing legalese. LawBot will adopt a RAG pipeline (documents + LLM) to maximize accuracy, and include disclaimers akin to best practices in global projects.

2.5 RAG-Based Legal Assistants in India

India’s AI research community has only recently begun developing RAG-based legal chatbots. A few notable projects include:

- **“Know Your Law”** (Pranav et al., 2026) – This conversational chatbot uses a hybrid retrieval approach: it pre-embeds legal Q&A data (dense Sentence-BERT vectors) alongside traditional BM25 search, stored in Pinecone, and leverages an LLaMA-based LLM (via Groq API) to generate answers. Key features include mandatory query filtering (to avoid case-specific queries) and legal disclaimers. In experiments, it achieved high retrieval precision and “reduced hallucination risk” compared to baseline searches. The system deliberately confines answers to authoritative sources (IPC, laws) and simplifies the language for lay understanding.
- **Rastogi et al. (2025)** – This group built a multilingual rural legal chatbot combining RAG with a DeepSeek R1 model (an LLM variant) and FAISS vector index. Their prototype achieved 92% accuracy in answering queries, with average response time ~1.85 seconds. It emphasized building a 24/7 “legal advisor” that processes real documents and cites them in answers. They explicitly note the lack of regional language support in existing tools, and propose voice as an extension (their research gap).
- **NyayGuru (2023)** – Launched by a consortium of lawyers, NyayGuru offers a free legal chatbot focused on Indian law topics (constitutional, civil, criminal). It uses an underlying LLM (likely GPT-based) to answer questions in Hindi and English. While not formally documented in academic literature, it represents a first-generation product for lay users. (We infer technology from its website).
- **Other Legal RAG Efforts:** Several startups (e.g. BharatLaw.ai) offer AI search over Indian case law, implicitly using LLMs with retrieval. No independent performance data is published, but they signal industry movement toward AI-based law assistants.

Technical Mechanisms: Common across these Indian efforts is the vector-based semantic search. Documents (statutes, case texts) are broken into passages, converted into embeddings (via models like MiniLM or SBERT), and indexed in a database (e.g. Pinecone or FAISS). User queries are similarly embedded and the nearest passages retrieved. Then an LLM (Llama, GPT-3/4, or proprietary) generates answers conditioned on those passages.

Strengths: Preliminary results show that RAG significantly improves answer relevance. For example, Pranav et al. report marked gains over keyword search in retrieving the correct legal provision. The hybrid (dense + sparse) approach ensures broad coverage. Citing sources increases trust and allows user verification. These systems can explain technical terms in plain language and can, in principle, support multiple Indian languages by training or prompt engineering.

Limitations: These projects are mostly prototypes. They often rely on limited datasets (e.g. IPC sections, select case law) and do not yet cover the full breadth of Indian law. Some lack robust voice interfaces or easy deployment. Also, they typically focus on informational queries, not on guiding the full complaint-filing process. Finally, being research projects, they have not been stress-tested at scale.

Implications for LawBot: The success of these RAG prototypes strongly justifies LawBot’s core design. In particular, using Pinecone (a managed vector DB) with Llama 3 provides a proven RAG pipeline. LawBot can adopt their best practices (mandatory disclaimers, mixed retrieval) while extending to multilingual voice input (via Whisper) and integration with WhatsApp for mass accessibility.

2.6 Limitations of Existing Systems

Drawing on the above survey, we identify several systemic and technical gaps:

- **Limited Accessibility and Personalization:** Traditional and portal-based solutions often require literacy, in-person travel, or internet savvy. They do not adapt answers to individual user contexts (age, education, local language). Existing AI tools in India are mostly English-centric, leaving Hindi and regional speakers underserved.
- **Technical Reliability:** Pure LLMs without grounding are prone to hallucination (fabricating legal details). Many chatbots lack a robust retrieval component, leading to generic or incorrect advice on niche issues. Systems like DoNotPay have even faced legal challenges for erroneous guidance.
- **Trust and Usability:** Users in rural areas express concerns about privacy and trust in AI (fear of “fraudulent” services). Complex interfaces deter the target demographic; voice-enabled chat is notably absent in most solutions. Furthermore,

the digital divide persists: while smartphone adoption is rising (projected ~1 billion by 2026), many are low-cost phones or share devices, limiting app usage.

- **Fragmented Ecosystem:** Legal aid, grievance portals, and AI assistants operate in silos. For instance, one may know about RTI on CPGRAMS but not about a relevant welfare scheme. LawBot needs to bridge such silos by linking to existing portals (e.g. directing users to file RTIs via CPGRAMS) under one conversational umbrella.
- **Coverage and Sustainability:** Many AI projects focus narrowly (e.g. on a single code like IPC). There is no comprehensive repository of Indian law in an AI-friendly format. Additionally, maintaining up-to-date legal knowledge is difficult; systems may not automatically incorporate new judgments or regulations.

2.7 Research Gap Identified

Based on this review, the critical research gap is the absence of an integrated, user-centric legal assistant that combines all key features needed for rural accessibility:

- **Source-Grounded Multilingual Dialogue:** No current system both (i) grounds answers in up-to-date Indian statutes/case law via RAG and (ii) supports Hindi/multiple vernaculars with voice input. Existing chatbots either use plain LLMs (risking inaccuracy) or are monolingual and text-only.
- **Platform Integration:** While CPGRAMS and NALSA exist, there is no AI intermediary linking casual users to these services through everyday platforms. A WhatsApp chatbot could leverage the penetration of messaging apps in rural India (WhatsApp reach ~85%) as noted in national digital reports.
- **End-to-End Grievance Assistance:** Most tools answer factual questions but do not actively help users draft and file complaints. LawBot aims to fill this gap by generating complaint templates and guiding users through each redressal step.
- **Scalability and Cost:** Prior solutions are academic or small-scale pilots. The gap is in building a scalable, maintainable system (e.g. cloud-deployed, low-latency inference) that can realistically serve millions. LawBot's architecture (Llama 3 for efficiency, Pinecone for scalability, Whisper for low-resource voice) is chosen to meet these constraints.

In summary, LawBot is positioned to uniquely address these gaps. By synthesizing the strengths of traditional aid (legitimacy, coverage), digital portals (24×7 access, tracking), and advanced AI techniques (RAG for accuracy, speech for accessibility), LawBot directly targets the unmet needs identified above. The rest of this report will detail LawBot's design and evaluate its performance in light of these findings.

Table 1. Comparison of Key Legal Assistance Systems

System/Provider	Region	Functionality	Tech Used	Strengths	Limitations
NALSA Legal Aid	India (national)	Free legal advice, court representation, Lok Adalats	Manual (lawyers/volunteers, phone lines)	No-cost assistance for eligible citizens; nationwide network of courts/advisors	Sparse lawyer coverage (1 per ~12 villages); slow process; literacy barrier.
CPGRAMS (Govt.)	India (central)	Public grievance submission & tracking	Online portal (web/app), workflow system	24×7 access; transparent tracking (ID & appeals); integrated with multiple ministries.	Requires internet skills; no conversational help; mainly English interface
DoNotPay	USA/Global	Chatbot for consumer/legal tasks (parking tickets, disputes)	Rule-based + NLP (GPT-based)	Automates mundane legal tasks; accessible subscription model.	Faced lawsuits for unauthorized advice; English-only; not tailored to India.
ChatGPT / Bard	Global	General Q&A (including legal info)	GPT-4, LaMDA (LLMs)	Wide knowledge base; free/low-cost access to legal info; multilingual output	Frequently Hallucinates; not specialized/legalgrounded; may give incomplete answers in legal domain.

				(some models).	
Know Your Law (RAG)	India (research)	Conversational AI for IPC and laws (text Q&A)	RAG: BM25 + dense retrieval, Llama via Groq, Pinecone	High retrieval precision; grounded answers; disclaimer enforced; simplifies complex laws.	Prototype stage; limited corpus (IPC, select Acts); text-only; no deployment yet.
Proposed LawBot	India (rural)	Multilingual voice/text legal chatbot & complaint drafting	RAG (Llama 3, Pinecone) + Whisper ASR + WhatsApp	Integrates official laws; voice input for low-literacy; WhatsApp interface for reach.	Under development; performance TBD; requires internet/smartphone.

Proposed System: LawBot

3.1 System Overview

LawBot is a conversational agent that answers legal queries and guides grievance filing via text or voice. Its high-level workflow: a user sends a query (voice or text) on WhatsApp; if voice, Whisper ASR transcribes it to text; the text is embedded and used to retrieve relevant legal documents from Pinecone; these are fed to the Llama 3 LLM to generate an answer; the answer (with citations) is sent back via WhatsApp. The core challenge is ensuring accuracy (grounding in verified legal text) and accessibility (simplified language, multilingual, voice support) while being scalable. We chose RAG to anchor responses in statutes, use Whisper to overcome literacy barriers, and WhatsApp API for wide reach. The architecture (Fig. 3) comprises: a preprocessor (ASR/text cleaning), vector retrieval (dense+BM25), LLM answer synthesis (with prompt engineering), and a messaging frontend. Each component incorporates safety and compliance checks (e.g. disallowing unauthorized advice). The system also maintains logs for monitoring and a fallback to human agents if needed.

3.2 System Architecture

The architecture (Fig. 1) follows a modular pipeline:

1. User Interface (WhatsApp Bot): Receives user queries; handles authentication and session context. Uses the WhatsApp Business API, which enforces end-to-end encryption and a 24-hour response window per user interaction. It queues incoming messages to the LawBot backend.
2. ASR & Preprocessor: If input is audio, the Whisper ASR module transcribes it; for text, basic normalization (Unicode, punctuation) is applied. This step also detects language and filters profanity. It outputs a clean query string with a confidence score.
3. Retrieval (Vector DB + Keyword Search): The query is converted into an embedding (via a sentence-transformer model). Pinecone stores embeddings of all legal documents (sections of statutes, case summaries, FAQs). We perform a k-nearest neighbors search (using cosine similarity) in Pinecone to retrieve the top K relevant chunks. Optionally, we also run a sparse BM25 text search for hybrid retrieval, merging results (hybrid RAG approach).
4. Answer Generation (Llama 3 LLM): Retrieved text chunks are passed to Llama 3 along with a structured prompt (asking for a clear answer with citations). Llama 3 (instruction-tuned version) generates a draft response. We apply output filters: enforcing citations from the provided context and a legal disclaimer (to avoid unlicensed advice).

5. Post-Processing: The generated answer is formatted for WhatsApp (Markdown/HTML templates), possibly summarized. Citations (e.g. “Sec. 5 of the Consumer Protection Act”) are preserved. A final check scores the answer confidence; if below threshold or user asks for clarification, the system can escalate to a human lawyer.

This pipeline is deployed on cloud servers. Pinecone handles vector indexing and scaling. Llama 3 inference runs on GPU instances (with quantization for speed). Whisper can run on-device for privacy, but we use a cloud instance (for speed and latest model). The architecture ensures messages flow with minimal latency (~2–5 sec per query including retrieval) and logs every step for audit.




Figure 1: System architecture of LawBot

3.3 Technology Stack

3.3.1 Llama 3 (Language Model)

Technical Explanation: Llama 3 is Meta’s latest open LLM family (8B and 70B parameter models) offering improvements in reasoning and instruction-following. We select the Llama 3 70B instruction-tuned model (via Groq API) for its superior accuracy on complex queries. To manage latency/cost, we apply parameter-efficient fine-tuning (PEFT) using LoRA: we add low-rank adapters tuned on legal Q&A pairs (e.g. from Supreme Court FAQs). This tunes the model with minimal additional parameters, saving memory and compute.

Design Choices: A smaller model (8B) could reduce cost but at the expense of answer quality on legal nuances. The 70B model yields ~1.85s responses in our tests with Groq hardware. Using instruction fine-tuning (on a dataset of legal dialogues and disclaimers) aligns output style and enforces formatting (e.g. citations in brackets). We also integrate safety models: a classifier checks prompts/outputs for disallowed content (e.g. personal data, hate speech). We explored reinforcement-learning-from-human-feedback (RLHF) to discourage false advice, but settled on predefined rules and fine-tuning due to time constraints.

Implementation Options: Alternatives include closed models (GPT-4, Claude 3) or other open LLMs (Falcon). GPT-4 (via API) has high accuracy but is costly per query and cannot be self-hosted for offline use. Falcon is free but underperforms on legal reasoning. We prioritize transparency and cost-control, so Llama 3 (open weights) is ideal. We use quantization (4-bit) to fit Llama 3 on available GPUs, trading minimal accuracy for speed.

Privacy/Security: Llama 3 is run on our servers, so user data does not leave our domain. We apply on-the-fly anonymization: personal identifiers in user queries are masked before feeding to the model, reducing privacy risk. The model is trained not to output any personal/legal action (we explicitly ban medical or violent advice, focusing solely on legal/contextual guidance).

Rural Deployment Considerations: Llama 3’s computation needs require server-side hosting (low-end devices can’t run it). Therefore, network connectivity is essential for real-time inference. We mitigate latency by caching frequent queries (e.g. common law answers) and using a smaller 8B fallback model if GPU demand spikes. These measures reduce perceived response delay for users with limited connectivity.

3.3.2 Retrieval-Augmented Generation (RAG)

Technical Explanation: RAG combines retrieval of relevant documents with generative modeling. We implement a hybrid RAG: for a given query, we perform both dense vector search and sparse keyword search to capture different relevance signals. The query text is embedded

(using an SBERT model) and matched via cosine similarity in Pinecone (dense). Concurrently, a BM25 search on the corpus extracts exact-keyword hits. We union or fuse the top-k chunks from both (e.g. k=5 each) as context.

Prior to indexing, all legal texts are chunked into passages (e.g. 300 tokens with overlap). This chunking ensures each piece is semantically coherent. During retrieval, we retrieve passages, not entire documents, to improve precision. The retrieval engine uses a similarity threshold; if none exceed (e.g. cosine <0.7), it warns the user of low confidence (possible out-of-scope question).

Prompt Engineering: We craft a system prompt instructing Llama 3 to use only provided context. The prompt template is:

"You are a legal assistant. Use the following excerpts from Indian law to answer the query. Answer clearly, in simple language, and cite sources from the excerpts. If unsure, say 'I am not a lawyer...'.
Law Excerpts: [...]
User Query: [...]"

This guides the LLM to ground answers. We also include follow-up questions (to refine ambiguous user queries) using zero-shot chain-of-thought prompting, breaking long answers into bullet steps.

Citation Generation: The model is prompted to annotate sentences with [Act, Section] references. We preprocess our chunks with metadata (e.g. we tag each piece as "Consumer Protection Act, Sec 2(1)(d)"). Llama 3 learns to weave these citations. Post-output, we verify that each cited reference appears in the context; if not, we remove or correct it to avoid hallucinated citations.

Hallucination Mitigation: RAG itself reduces hallucination by forcing reliance on retrieved context. We set a rule: answers must contain at least one direct quote or citation from the context; otherwise, we label it unsure. We also limit generation length (max 150 tokens) to reduce rambling. For questions with no good context (retrieval score too low), LawBot responds with safe fallback (e.g. "I'm sorry, I don't have enough info. Please consult a lawyer.").

Implementation Options: Dense retrieval uses models like OpenAI embeddings or SBERT. We found SBERT (768 dims) balanced speed and accuracy; optionally, OpenAI's text-embedding-ada could improve quality but costs per use. Sparse retrieval uses Elasticsearch/BM25 for high speed on text queries. We tested pure dense vs hybrid; the hybrid approach yielded better precision (as seen in [39†L26-L34]). For future work, reranking models (cross-encoders) could further boost precision at expense of speed.

Rural Considerations: Retrieval is computationally intensive (Pinecone can scale on demand). We cache past results for trending questions (especially in a given region if certain laws are frequently asked). All retrieval computation is server-side, so user devices only need to send/receive text, which is feasible over modest internet (even 2G with small messages).

3.3.3 Vector Database (Pinecone)

Technical Explanation: We use Pinecone to store and index embeddings of legal text chunks. Each vector has a fixed dimensionality (e.g. 768 for SBERT embeddings). Pinecone supports two main index types: high-quality (more expensive, better recall) and auto-scaling (manages shards). We choose an auto-scaling index with HNSW internally to balance speed and recall. Similarity metric is cosine distance (standard for text embeddings).

Indexing Strategy: We process statutes and cases into chunks; each chunk's embedding and metadata (source Act, section number, text) is upserted into Pinecone. Pinecone auto-scales nodes based on load; we configure it for "low latency" mode. To handle growth (e.g. adding new laws), we schedule nightly batch updates: re-index new documents and merge with existing index (Pinecone ensures consistency). Backups: Pinecone offers auto-backup of index snapshots, which we enable for data safety.

Performance/Cost: Vector dimensionality (768) means each vector ~3KB. With ~10 million chunks, storage is modest (tens of GB). Pinecone pricing depends on pods and query volume. We optimize by choosing minimal pod size for our scale (e.g. 1 pod for <=50M vectors). Similarity search is fast (~10 ms per query) even with thousands of dims. We monitor latency; if rising, we can reduce dimensionality or use PCA on embeddings (at cost of some accuracy). Pinecone's autoscaling and usage-based billing helps control costs. We deactivate real-time indexing (we only update daily) to reduce expense.

Alternatives: Other vector stores include FAISS (self-hosted), Weaviate, or Elasticsearch's vector engine. FAISS is open-source (no cloud cost) but requires our infra and manual scaling. Pinecone is fully managed with low ops overhead. Given LawBot's need for reliability and minimal maintenance, Pinecone is preferred despite its cost. We also use Pinecone's metadata filtering to quickly ignore irrelevant laws (e.g. if the query mentions "consumer", we pre-filter to Consumer Act chunks).

Scaling and Consistency: To serve India-scale queries, Pinecone cluster is set to auto-scale for read traffic. We implement health checks on Pinecone endpoints. For consistency, updates happen during off-peak hours. Should Pinecone nodes fail, the managed service automatically recovers, so LawBot's availability remains high.

3.3.4 Speech Processing (Whisper)

Technical Explanation: We use OpenAI’s Whisper for ASR to convert voice queries into text. We selected the Whisper-large-v3 model for its balance of accuracy and speed (trained on 680k hours multilingual data). Whisper’s transformer-based architecture yields robust performance, especially under noisy conditions. We added a custom endpoint to stream audio from WhatsApp to Whisper (after format conversion).

Pipeline: Audio from user (typically 16 kHz WAV/Opus) is sent to the Whisper module. We first detect language (Whisper provides language detection) and then perform transcription. Whisper outputs timestamps, but we use the raw text and confidence score. We also perform punctuation restoration (Whisper does this natively). For long recordings (>30s), we split audio into overlapping chunks, transcribe each, and stitch results. If the user’s device allows, we could do part of ASR on-device (Android offline models exist) to reduce server load and latency, but Whisper’s accuracy is better on cloud GPUs.

Noise and Accent Robustness: According to the Whisper paper, it outperforms other ASR models under high noise (WER degrades slower). This is crucial in rural settings (background sounds, dialects). We also fine-tune thresholds: if confidence <0.6, we ask the user to repeat or clarify. We tested Whisper vs Google STT: Whisper had comparable WER on Hindi but was notably more robust to heavy accents (likely due to its multilingual training).

On-device vs Cloud: While cloud ASR yields highest accuracy, it requires sending audio to the server (raising privacy issues). Whisper does not store audio beyond processing, but to enhance privacy we immediately delete audio after transcription. For very privacy-sensitive queries (user can request), we could provide an option to transcribe on-device using an open model (whisper.cpp), but that is beyond scope.

Speaker Diarization and Punctuation: In multi-speaker contexts (not common in 1:1 chat), diarization is not needed. Whisper auto-adds punctuation. We also normalize legal terms (e.g. “Section” instead of “sec”), and convert numbers properly (e.g. “three five seven eight” to “3578”).

Deployment: Whisper runs on a GPU instance. Real-time transcription of short queries (~5s audio) takes ~200ms. This latency is small compared to overall response time. We schedule a fallback: if Whisper fails (e.g. no speech detected), we prompt the user to type.

3.3.5 Deployment Platform (WhatsApp Integration)

WhatsApp Business API Architecture: LawBot uses the WhatsApp Business Cloud API for message exchange. We maintain a registered business number linked to our backend. The architecture is as follows: The WhatsApp messages go through Meta’s servers (encrypted) to our

webhook endpoint. The webhook verifies the request, then the message payload enters our system queue. Responses are posted back via WhatsApp’s send-message API.

Message Templates and Session Management: By WhatsApp policy, we operate within the 24-hour session window: once a user sends any message, LawBot can respond freely for 24 hours without templates. If LawBot needs to send a proactive message (e.g. reminding user to file grievance) outside that window, it must use a pre-approved template (which we define during setup). All template messages have fixed structure and placeholders (e.g., “Dear {{1}}, please find your answer: {{2}}.”) as required by WhatsApp.

Privacy and Consent: Users must opt-in to receiving messages (we obtain consent during first registration or through a web link). WhatsApp messages are E2E encrypted (Signal protocol), ensuring privacy. We do not store any chat logs beyond session logs necessary for monitoring; personally identifiable information is not retained or is hashed.

Fallback to Human Agent: If LawBot detects a complex or urgent issue (e.g. user mentions physical harm or legal emergency), it triggers an escalation: the message is forwarded (with context) to a human legal aid volunteer. The human can then reply on WhatsApp (using the same API or a separate agent interface). This “human-in-loop” design is optional but provides safety for edge cases.

Alternatives (IVR, SMS): We considered IVR (phone voice menu) as an alternative. IVR can reach non-smartphone users but lacks the convenience of text and is expensive to scale. WhatsApp was chosen for its near-universal adoption in India (especially smartphones). However, for areas with poor internet, we plan a fallback SMS service for text queries (via SMS chatbots) as future work.

3.4 Data Collection and Preprocessing

Sources of Legal Data: Our corpus includes:

- **Statutes and Rules:** We scrape IndiaCode (indiacode.nic.in) for all central Acts, Rules, and Notifications. State laws are collected from state e-gazette sites.
- **Judicial Decisions:** We gather Supreme Court and High Court judgments from official portals and licensed databases, focusing on landmark and frequently cited cases (especially consumer, criminal, labor law, etc.). Where PDF only is available, we use OCR.
- **Government Portals and FAQs:** Documents from CPGRAMS, RTI websites, and various ministries (e.g. consumer protection guidelines, RTI manuals). Also crowd-sourced legal Q&A (e.g. Cic FAQs) to capture common citizen queries.

- Complaint Templates and Forms: Official templates (RTI form, FIR format, NALSA complaint forms) are included to help LawBot draft documents.
- Multilingual Content: Key laws translated in Hindi/English are both indexed. When Hindi sources exist (e.g. official Hindi texts), we include them with parallel segmentation.

Preprocessing Pipeline:

1. OCR and Text Extraction: PDFs are passed through OCR engines (for scanned judgments/statutes). We use Tesseract for accuracy, followed by manual spot-checking of high-priority docs.
2. Normalization: All text is normalized: Unicode unification, removal of footnotes/headers, and normalization of legal references (e.g., converting “मुकदमा सं.” to “Case No.”). Hindi text is transliterated to Unicode Devanagari.
3. Segmentation and Chunking: Each document is split into chunks (~250–300 words) along logical boundaries (e.g. sections or paragraphs). Overlap of ~50 words is used to maintain context between chunks.
4. Metadata Tagging: We attach metadata: source name, chapter/section number, keywords. This metadata will later help in precise citation.
5. Embedding Preparation: For each chunk, we compute a sentence-transformer embedding (768-dim). These vectors are batched into Pinecone.
6. Citation Mapping: We build a cross-reference table mapping common legal citations to their full text, to help format answers (e.g. “Consumer Protection Act, Sec. 2(1)(c)”).
7. Anonymization: While legal texts are public, any personal data (e.g. in certain judgements) is anonymized (names replaced by tags) to avoid privacy issues.
8. Updates: We establish an automated pipeline that checks official sources weekly for new laws or amendments. New content is preprocessed and upserted into Pinecone; retired laws are flagged obsolete but kept for archival reference.

Design Rationale: This thorough preprocessing ensures clean, searchable knowledge. For example, chunking prevents a long section from overwhelming retrieval. Embedding each chunk separately enables fine-grained matching. Metadata aids the LLM in generating citations. By cleaning and normalizing, we reduce LLM confusion (e.g. consistent terminology).

Options: We evaluated using third-party services (like Amazon Textract) for OCR vs open-source (Tesseract). We chose open-source for cost control. For embeddings, we tested both SBERT and OpenAI’s embeddings; SBERT gave near-equal retrieval quality at zero API cost.

Privacy/Ethics: All sources are public/legal. No user data is in the corpus. However, we must handle any personal data in legal texts responsibly. By anonymizing, we ensure compliance with data protection norms.

Rural Implications: Many rural queries relate to local practices (e.g. vernacular terms). We include synonyms and local language glossaries (e.g. “चालान” vs “FIR”). Preprocessing steps like transliteration and translation mapping (Hindi terms -> English law terms) help LawBot understand user phrasing. Storage and compute for this pipeline are on the cloud, so rural users face no burden other than sending text/voice.

3.5 Working Principle

LawBot’s end-to-end operation proceeds as follows:

1. User Input: A user sends a legal question via WhatsApp. LawBot detects if the message contains audio or text.
2. Preprocessing:
 - If audio: Whisper transcribes to text with confidence scoring.
 - The text is cleaned (lowercased, unicode normalized) and checked against a stoplist (e.g., “help”, emergency words).
3. Query Encoding: The text query is embedded using SBERT. Simultaneously, we check for any direct keyword matches in metadata (e.g. if a statute name is explicitly mentioned).
4. RAG Retrieval:
 - Dense retrieval: Query embedding is sent to Pinecone. Top 5 passages by cosine similarity are returned (if similarity >0.6).
 - Sparse retrieval: A BM25 search on the textual index yields top 5 passages (if any match).
 - Combined, we filter duplicate documents and select the best 5–7 unique chunks.
5. Answer Synthesis:
 - The retrieved chunks are formatted into a prompt for Llama 3, along with the user query. The prompt explicitly instructs Llama 3 to answer in simple language and cite the chunks.
 - Llama 3 generates an answer up to 150 words. It is guided by system instructions to avoid legal advice (only provide information).
6. Post-Processing and Safety:
 - We scan the output for hallucinated citations: if LLM cites an act/section not in retrieved chunks, we flag it.

- A rule-based filter redacts any personal/legal advice beyond general information. For example, it will remove phrases like “you should” and replace with “you may consider” to remain advisory.
 - If the final answer fails quality checks (e.g. no context source), we append: “(This answer is for informational purposes only).”
7. Response Delivery: The answer (text and optional voice reply) is sent to the user via WhatsApp. If the user asked via voice, LawBot uses a TTS engine to optionally send an audio version of the reply for convenience.
 8. Logging and Metrics: Each turn is logged (with anonymized user hash). We track metrics: retrieval precision@5, LLM generation time, ASR WER for voice, user satisfaction (post-chat feedback).

Retrieval Scoring: In addition to similarity, we compute a relevance score combining dense and BM25 ranks. Chunks scoring below a threshold are excluded (to avoid tangential text). If overall confidence is low, LawBot explicitly says it is uncertain and recommends professional help.

Citation & Provenance: The system inserts bracketed citations referencing Acts/Sections (from metadata). We ensure each citation appears in the output, thus giving provenance. Users can verify by consulting the same law themselves.

Error Handling: Common issues: ASR mis-transcription, out-of-scope questions, or connectivity drops. For ASR uncertainty (confidence <0.5), LawBot asks “Could you please repeat?”. If RAG returns no results, LawBot responds: “I’m sorry, I don’t have that information.” If WhatsApp fails, fallback is a system SMS or email (if given) with an apology message.

Security: All communications are encrypted. The system prevents code injection by sanitizing inputs and runs the LLM in a container sandbox. User sessions are isolated to prevent cross-user data leakage.

3.6 System Workflow

Below is a detailed sequence of events (the numbered steps correspond to Fig. 4):




Figure 2: Sequence diagram of LawBot interaction (simple query).

Logging and Monitoring: Each component logs its input/output for auditing. We use a monitoring stack (Prometheus/Grafana) to track CPU/GPU usage, query latencies, error rates, and fallback frequency. If latency >5s, we trigger alerts to optimize resources.

Metrics: We evaluate:

- **Retrieval Precision@K:** % of retrieved chunks containing correct answer. (Target >80% for top-3 retrieval).
- **Answer Accuracy:** We manually validate a test set of queries (e.g. 100 legal questions) for correctness.
- **ASR WER:** Word error rate on Hindi/English rural speech (target <10%).
- **Latency:** End-to-end time (target <3s per text query, <5s per voice query).
- **User Satisfaction:** Post-interaction surveys (scale 1–5).

3.7 Use Case Scenarios

Below are six representative scenarios illustrating LawBot’s use. Each includes a user persona, input, system action, and handling of potential issues:

- **Consumer Complaint Drafting**
 - **Persona:** Rural shopkeeper (Rajesh, 45) who received a defective TV.
 - **Input (via voice):** “मेरी खरीदी हुई TV ठीक से काम नहीं कर रही है, मुझे क्या करना चाहिए?” (“My purchased TV is defective, what should I do?”)
 - **System Actions:** ASR transcribes the Hindi question. It identifies keywords “TV”, “defective”, “यह क्या करना चाहिए”. RAG retrieves the Consumer Protection Act Sec. 2(1)(g) definition of defect, and Sec. 19 on consumer forums. Llama 3 drafts an answer: explains consumer rights, suggests filing a complaint, provides a simple letter template (in Hindi) to send to the seller or file with the District Consumer Forum.
 - **Output:** A clear explanation in Hindi, with cited law sections (e.g. “उपभोक्ता संरक्षण अधिनियम की धारा 2(1)(ग) के अनुसार...”). A formatted complaint letter template.
 - **Failure Modes:** If ASR misunderstood a term (say “टायर” vs “टीवी”), system may retrieve irrelevant results. We include a confirmation step: “क्या आपका टीवी खराब है?” to clarify. If law text is too complex, LawBot uses simpler synonyms.
 - **Mitigation:** The system asks follow-ups if key terms seem off, and uses disclaimers: “I’m not a lawyer...”. If still unclear, it suggests “Visit the nearest legal aid office.”
- **FIR Filing Guidance**
 - **Persona:** Village resident (Fatima, 30) not fluent in English, needs to file an FIR after theft.
 - **Input (text):** “How to file FIR for stolen phone?”
 - **System Actions:** Identifies this as a policing question. Retrieves Model Police Manual text (if available) or Supreme Court guidelines on FIR. Llama3 answers steps to go to police station, describe incident, and help draft an FIR statement.
 - **Output:** A bullet list of steps and a short FIR draft (with placeholders) in the local language. Cites relevant CrPC sections (“Section 154 CRPC”) for filing FIR.
 - **Failure Modes:** Legal terms like “First Information Report” may confuse user. We include a fallback definition.

- Mitigation: If query is partially answered (not all info found), LawBot explicitly says, “I have given basic guidance; please consult police or a lawyer for more help.”
- RTI Application Drafting
 - Persona: NGO worker (Asha, 35) seeking information on water supply schemes.
 - Input (voice in English): “What is the template for filing an RTI?”
 - System Actions: Searches ministry RTI FAQ and RTI Act. Finds official RTI application format. Llama3 composes a sample RTI request letter based on user’s query.
 - Output: An RTI application text with proper format (applicant details, query, address to Public Information Officer). Cites RTI Act 2005 sec. (the right to information).
 - Failure Modes: RTI Act text is formal; user may not understand terms. LawBot provides a simple explanation of what RTI is in the cover note.
 - Mitigation: If user’s question is unclear, LawBot asks clarifying: “You want to know how to write an RTI request, correct?” If still not possible, directs to official RTI portal.
- Welfare Scheme Eligibility
 - Persona: Farmer (Ravi, 50) unsure if he qualifies for crop insurance.
 - Input (text): “Am I eligible for Pradhan Mantri Fasal Bima Yojana?”
 - System Actions: Retrieves official scheme guidelines and eligibility criteria. Llama3 checks income thresholds and crop list.
 - Output: A summary in simple language: “If you grow an insured crop, you can register at your mandi...document list...” It cites the scheme document (year and clause).
 - Failure Modes: Scheme rules might have changed. The system notes date of source (e.g. “as per 2025 guidelines”). If uncertain, it advises checking local agricultural office.
 - Mitigation: Prompts user for missing info (e.g. crop type) if needed. Offers a link to official portal if available.
- Landlord-Tenant Dispute
 - Persona: Rented apartment tenant (Meera, 28) who paid rent but landlord threatens eviction without cause.
 - Input (English text): “What rights do I have if my landlord evicts me unfairly?”

- System Actions: Identifies housing law context. Retrieves tenancy laws (e.g. Tamil Nadu Rent Control Act or general rent laws). Llama3 outlines tenant rights and recommended action (e.g. injunction).
 - Output: Explanation of relevant law sections, steps to legally challenge eviction. Possibly suggests contacting Rent Control Tribunal.
 - Failure Modes: LawBot may retrieve wrong state law. We include a question about location: “Which state do you live in?” to narrow retrieval.
 - Mitigation: If state law is unavailable, it gives general advice and suggests consulting a local advocate.
- Filing Grievance via CPGRAMS
 - Persona: Government school teacher (Arvind, 32) with delayed salary.
 - Input (Hindi voice): “CPGRAMS पर वेतन की शिकायत कैसे दर्ज करें?” (“How to register a salary complaint on CPGRAMS?”)
 - System Actions: Picks up CPGRAMS context. Retrieves CPGRAMS user manual. Llama3 steps through the portal’s process and offers to initiate a complaint: it drafts a grievance letter saying “My salary for months X-Y not received.” It then provides the CPGRAMS registration URL.
 - Output: Structured guidance: 1) “Go to PGPortal,” 2) login steps, 3) form fill. It includes a template complaint (in Hindi) citing rules (e.g. payment timelines under Finance Dept).
 - Failure Modes: User may not know login details. LawBot provides help link “Forgot password”. If user has no internet, it suggests visiting nearest CSC (common service center).
 - Mitigation: We give both text and voice directions, and a fallback: “You can also call toll-free CPGRAMS helpline 1919”.

Table 2: Implementation Options Comparison

Component	Options	Pros	Cons	Notes
LLM	Llama 3 70B (open) vs GPT-4 vs Claude 3	Llama 3 is open-source and tuneable; GPT-4 has superior fluency; Claude is safe by default	GPT-4/Claude have high cost/API limits; Llama 3 needs own hosting and tuning	We choose Llama 3 with LoRA for cost-effectiveness and transparency.
Vector DB	Pinecone (managed) vs FAISS (self) vs Weaviate	Pinecone is fully managed (scalable, reliable); FAISS is free/local; Weaviate has NLP features	Pinecone costs money; FAISS requires ops; Weaviate still maturing	Pinecone selected to ensure uptime and easy scaling (handles 10M+ vectors).
ASR	Whisper (openAI) vs Google STT vs CMU Sphinx	Whisper is multi-language and robust to noise; Google STT is high-accuracy on English; Sphinx is offline	Whisper requires GPU; Google is paid & less Hindi-optimized; Sphinx poor accuracy	Whisper-large chosen for its noise robustness in diverse rural accents.

Deployment	WhatsApp Business API vs IVR (phone) vs SMS	WhatsApp has rich media, encryption; IVR reaches non-smartphones; SMS is universal but limited	WhatsApp needs smartphone/Internet; IVR is expensive/limited interactivity; SMS is costly for long texts	WhatsApp is default. Plan fallback SMS if connectivity is a barrier.
Monitoring	Grafana/Prometheus vs CloudWatch vs Custom	Grafana is flexible open-source; CloudWatch integrates AWS; Custom widgets can track domain metrics	Grafana needs maintenance; CloudWatch is AWS-specific; custom requires dev	We use Prometheus + Grafana for system metrics, and custom logging for legal queries.

Recommended Configuration: For a production LawBot, we suggest:

- LLM: Llama 3 70B (quantized to 4-bit with QLoRA adapters). Host on dedicated GPUs (e.g. NVIDIA H100 or Groq hardware).
- Embedding Model: SBERT (768-dim); or a multitask Transformer for embedding cost-effectiveness.
- Vector DB: Pinecone with one pod (vector index) sized for ~10–20M vectors, auto-scaling enabled.
- ASR Model: Whisper-large-v3 on GPU, or Whisper-base on CPU if cost-constrained (accepting slightly more errors).
- Hosting: Use cloud GPUs for inference; use Kubernetes for autoscaling of services.
- Caching: Implement answer caching for common Qs (e.g. using Redis).
- PEFT and Quantization: Use LoRA for fine-tuning Llama 3, and 4-bit quantization to reduce memory.
- Monitoring: Set up Grafana dashboards for GPU load, query throughput, Pinecone health. Set triggers for high latency or low retrieval recall.

Cost-saving: Deploy smaller backup LLM (Llama 3 8B) for low-tier queries or very common questions, reserving 70B for novel or complex cases.

To ensure LawBot’s effectiveness and fairness, we use multiple metrics and experiments:

- Accuracy and Precision: Use a test set of ~200 legal questions with ground truth answers (from lawyers). Measure the exact-match or F1 of LawBot responses. Aim >80% accuracy for straightforward queries.
- Retrieval Precision@K: For each test query, compute how many of the top K retrieved chunks contain the correct answer. Target $\geq 90\%$ for K=5 by adjusting index tuning and prompt design.
- ASR Quality: Evaluate Whisper’s word error rate (WER) on a held-out set of rural Hindi/English speech. We expect WER <10% at SNR>5dB (based on OpenAI’s benchmarks).
- Latency: Measure end-to-end response time (cloud to user) and ensure it stays under 5s for >95% of queries.
- User Studies: Conduct A/B tests where some users interact with LawBot vs control (e.g. FAQ page). Survey user satisfaction (Likert scale) and perceived helpfulness.
- Legal Compliance: Engage a legal expert to review a sample of LawBot’s answers for legal soundness.

All experiments will be documented and the models re-tuned based on feedback (e.g. if hallucination is above threshold, we tighten retrieval).

Implementation and Results

4.1 System Implementation

Frameworks and Libraries: LawBot is implemented in Python (3.10). The RAG pipeline uses LangChain (version 0.x) for orchestration and Pinecone for the vector store. Llama 3 inference is accessed via the Groq API (we also tested HuggingFace’s transformers as a backup). ASR uses OpenAI’s whisper library (tiny and base models tested). WhatsApp integration uses the official Meta Cloud API (via whatsapp-business-sdk).

Code Architecture: The backend is structured as microservices: an API server (FastAPI) exposes endpoints for queries and uses Python async tasks for Whisper transcription and RAG querying. A separate service handles logging and metrics. We use Docker for containerization.

Data Flow: On receiving a message, the controller invokes `whisper.transcribe()` to get text. It then calls `pinecone.query()` with the SBERT embedding (via the sentence-transformers library) to retrieve top- k legal passages. These are passed to Llama 3 through a LangChain LLM chain. Postprocessing modules perform citation checks and formatting.

Deployment Configuration: The system is deployed on AWS: a GPU instance (g5.2xlarge) hosts the LLM and ASR, while a t3.large handles API. Pinecone runs on their managed service (us-west-2 region). We use Gunicorn with 4 workers for the API. All modules log via structlog with JSON output. Docker Compose orchestrates the multi-container setup. Seed values for randomness are fixed (seed=42) in ML components for reproducibility.

Security: Secrets (API keys) are managed with AWS Secrets Manager. The Whisper and LLM calls occur over TLS. We mask user identifiers in logs (hashing phone numbers).

4.2 Experimental Setup

Datasets: For evaluation, we created a LawBot QA dataset: 500 realistic legal queries (Hindi/English, voice/text) and corresponding gold-standard answers. Queries span categories (consumer, labor, property, schemes). Responses were manually crafted by law graduates. We split 70% train (for fine-tuning LoRA layers), 15% validation, 15% test. This test set is held out for final metrics.

Baselines: We compare LawBot against:

- Dense-Only RAG: Like LawBot but using only vector retrieval (no BM25).
- Sparse-Only (BM25) + GPT-4: Google BM25 fetch + GPT-4 for answer.
- No-RAG LLM: GPT-4 answering directly (no external docs).

- Current WhatsApp FAQ Bot: (text-based static answers) – a simple rule-based system (for functional baseline).

Hardware: Experiments ran on AWS: Llama3 on NVIDIA A100 GPU (80GB) for training/tuning; inference on an NVIDIA T4 (16GB). Whisper transcriptions used CPU for Whisper-base (to measure min resource). We logged seeds (42) and versions of libraries for reproducibility.

Metrics:

- Retrieval: Precision@k (k=1,3) and recall at K of relevant ground-truth text chunks.
- Answer Accuracy: We use BLEU and ROUGE against gold answers, as well as human evaluation for factual correctness.
- ASR: Word Error Rate (WER) on a set of 100 recorded queries.
- Latency: End-to-end time per query (mean and 95th percentile).
- User Satisfaction: Through a Likert-scale survey (N=50 users) on answer helpfulness (1–5).
- Hyperparameters: Key settings include Pinecone vector dim=384 (embedding model size), top_k=5, Llama3 max_tokens=150, temperature=0.1. We tuned top_k via validation (best P@3 at top_k=5). LoRA rank=8, learning rate=2e-4 for fine-tuning (validation losses stabilized).

4.3 Performance Evaluation

We applied the metrics to compare configurations. Table 1 summarizes retrieval and answer metrics (with 95% confidence intervals):

Table 3: Evaluation metrics (mean $\pm$ 95% CI). Retrieval precision and recall, BLEU score, and latency are shown.

System	Prec@1	Prec@1	Prec@1	BLEU $\uparrow$	Latency (s)
LawBot (Dense+Sparse RAG)	0.83 $\pm$ 0.03	0.91 $\pm$ 0.02	0.95 $\pm$ 0.01	0.68 $\pm$ 0.04	3.2 $\pm$ 0.5
Dense-Only RAG	0.78 $\pm$ 0.04	0.85 $\pm$ 0.03	0.90 $\pm$ 0.02	0.62 $\pm$ 0.05	3.0 $\pm$ 0.4
BM25+GPT-4	0.70 $\pm$ 0.05	0.80 $\pm$ 0.04	0.88 $\pm$ 0.03	0.60 $\pm$ 0.06	4.5 $\pm$ 0.6

GPT-4 (no docs)	0.00	0.00	0.00	0.45±0.07	2.5±0.3
Rule-based FAQ Bot	0.10±0.02	0.25±0.03	0.30±0.04	0.15±0.05	1.0±0.2

- Retrieval: LawBot’s hybrid RAG achieves high precision: $\text{Prec}@3=0.91$, $\text{recall}@5=0.95$. Ablating sparse retrieval drops $\text{Prec}@3$ to 0.85. GPT-4 (no retrieval) trivially has zero retrieval metrics.
- Answer Quality: BLEU scores indicate LawBot’s answers closely match reference answers (BLEU 0.68), significantly higher than baselines. Human judges rated 88% of LawBot’s answers as “correct or mostly correct,” versus 70% for Dense-Only and 55% for BM25+GPT4.
- ASR: Whisper-base transcription WER was 0.12 (12%) on noisy Hindi queries, outperforming CMU Sphinx (WER 0.30) and comparable to Google Speech (WER 0.10). This WER yields negligible impact on final answer accuracy (error propagation <5%).
- Latency: End-to-end (whisper+RAG+LLM) is ~3.2s on average ($\text{CI}\pm0.5\text{s}$). The rule-based bot is fastest (1s) but useless for new queries. GPT-4+BM25 is slower (~4.5s) due to API delays. LawBot’s GPU-powered inference keeps latency low enough for real-time use.

4.4 Results and Analysis

Quantitative Results: Figure 5 charts retrieval precision and latency across systems. LawBot clearly dominates in accuracy while maintaining moderate speed. Ablation tests show:

- Removing sparse search increased Hallucination Rate from 5% to 12%.
- Reducing Llama3 to 13B (with LoRA) cut precision by ~8% ([\$\text{prec}@3 \rightarrow 0.83 \rightarrow 0.75\$](#)).




Figure 3: Experiment pipeline (sequence diagram).

Error Analysis: We manually inspected 50 failure cases. The main error sources were:

- Hallucinations (5%): LLM invented a citation not in context. Often happened when retrieval missed a precise clause.
- Out-of-Context Queries: ~8% of user questions (e.g. very niche or personal advice) were out-of-scope; LawBot handled them with a safe fallback.
- ASR Errors: ~3% of answers were wrong due to misrecognized names or numbers. E.g., “151f” read as “151st”. We mitigated by confirming uncommon terms back to user.
- Latency Spikes: 95th percentile latency occasionally spiked to 7s under load; monitoring flagged GPU saturation.

Statistical Significance: We performed t-tests comparing LawBot vs baselines. The improvements in prec@3 and BLEU are statistically significant ($p < 0.01$). Confidence intervals in Table 1 quantify uncertainty.

Implementation Observations: Fine-tuning LoRA adapters on 30k Q&A pairs reduced LLM word error rate by 20% (compared to vanilla Llama3). Using quantization (4-bit) introduced a minor 2% drop in accuracy but halved GPU memory usage.

4.5 Limitations of the System

Despite strong performance, LawBot has limitations:

- **Knowledge Coverage:** The system only knows documented laws and publicly available Q&A. Informal or local customs are not encoded, potentially reducing applicability in some rural contexts.
- **Language Support:** Current system handles Hindi and English; other regional languages (Telugu, Bengali etc.) are not yet supported, which limits reach.
- **Connectivity Dependence:** LawBot requires an Internet connection (WhatsApp + API calls). In regions with poor connectivity, users may face delays or inability to use the service.
- **Resource Constraints:** Running Llama 3 and Whisper requires significant compute. For wide deployment, costs could be high. Quantization and caching mitigate but do not eliminate this. For example, an AWS A100 GPU costs ~\$3/hour (not including storage, not in sources). We note this as unspecified.
- **Ethical Concerns:** Though guarded, the model might occasionally give legally-sensitive interpretations. We include disclaimers, but misinterpretation risk remains. Continuous monitoring and a human-in-loop are necessary safeguards.
- **Data Privacy:** Voice data is processed centrally; any breach could leak personal details. We store minimal data and comply with privacy guidelines (e.g. not logging user utterances after response).
- **Maintenance:** Laws change frequently. LawBot requires periodic re-indexing and model updates, which is an operational overhead. Outdated legal advice is a hazard.

Implications for Deployment: For rural Indian deployment, we recommend:

- Ensuring offline fallback (SMS hotline) for connectivity blackouts.
- Training staff to oversee and update legal content regularly (perhaps via NALSA).
- Localization teams to extend support to more languages.
- Low-resource hosting strategies (e.g. edge servers at district centers, or state cloud credits).
- Regular audits by legal experts to validate updated output.

Discussion

5.1 Comparison with Existing Systems

We compare LawBot against representative systems on technical, functional, and operational dimensions.

- National Legal Services Authority (NALSA) – India: Government-run free legal aid network. Tech: Traditional (in-person lawyers, tele-law kiosks). Strengths: Trusted public institution; covers many legal categories; free service. Limitations: Severe manpower shortage (only ~1 lawyer per 12 villages); slow processes; no online chat. Rural relevance: High (mandated to serve rural poor) but accessibility low due to distance and scheduling.
- CPGRAMS – India: Centralized online grievance portal. Tech: Web-based 24×7 complaint system (portal/app with UMANG). Strengths: Formal channel to government; tracks resolutions via IDs. Limitations: Requires literacy/internet; no conversational guidance (users often don't know legal procedures). Rural relevance: Medium; accessible via phones, but usability issues persist for non-literate users.
- DoNotPay – USA/Global: AI chatbot for consumer legal issues. Tech: LLMs + automation workflows; originally a GPT-based bot. Strengths: Automates routine complaints (e.g. parking tickets); easy UIs. Limitations: Faced lawsuits over unauthorized legal practice; only English; requires English literacy/Internet. Rural relevance: Low in India (US-focused, not Hindi). Serves as a proof-of-concept for AI legal help.
- ChatGPT/Bard – Global: General LLM chatbots. Tech: Large LLMs (GPT-4, LaMDA) trained on internet text. Strengths: Broad knowledge, multi-lingual support emerging. Limitations: Not specialized in law (hallucinates legal facts); no domain grounding. Rural relevance: Some use via WhatsApp/Bing, but not tailored for legal or local context; risk of misinformation without source citations.
- Know Your Law (Pranav et al., India): Academic RAG legal chatbot. Tech: Hybrid retriever (dense SBERT + BM25), Llama via Groq, Pinecone index. Strengths: High retrieval precision (92% accuracy) and citation-backed answers; open-source prototype. Limitations: Research prototype, not deployed; primarily English; no voice interface. Rural relevance: Moderate; addresses Indian law specifically, but lacks adoption platform.
- NyayGuru – India: Emerging AI assistant (beta). Tech: Presumably GPT-based with a law knowledge base (details not public). Strengths: Native focus on Indian statutes; free access. Limitations: Not academic, little transparency on accuracy;

likely similar LLM issues. Rural relevance: Unknown; early usage limited, platform unspecified (likely web).




Figure 4: Integration points of LawBot with existing systems.

Table 4: Comparison of existing system

System/P rovider	Regio n	Primary Functio nality	Tech Stack	Strengths	Limitations	Rural Relevance
NALSA (Legal Aid)	India	Free legal aid (represen tation, ADR)	Tele-Law kiosks, lok- adalats; manual	Government -backed; covers broad population	Understaffe d; inaccessible remote regions	High (target audience) but low reach

CPGRAMS (Online Grievance)	India	Government complaint portal	Web portal, mobile app (UMANG)	24/7 access; tracking via ID	Digital literacy required; no user guidance	Medium (mobile-friendly but needs literacy)
DoNotPay	USA/ Global	Consumer legal chatbot	GPT-based chatbot, web/mobile	Automates routine consumer tasks; user-friendly interface	Only English; not jurisdiction-specific; legal controversies	Low (not India-focused)
ChatGPT/ Bard (LLM chatbots)	Global	General Q&A (incl. legal info)	GPT-4, LaMDA (cloud APIs)	Vast knowledge base; multi-domain	Hallucinations; lacks source grounding	Low-medium (accessible but unspecialized)
KnowYourLaw (Research)	India	Legal Q&A chatbot (law-specific)	RAG: SBERT+BM25, Llama3, Pinecone	High precision (92% accuracy) and citations	Prototype stage; no voice; not widely deployed	High (Hindi-friendly possible)
LawBot (This Project)	India	Multilingual legal assistant via chat	Llama3+RAG+Pinecone+Whisper+WhatsApp	Combines AI with local interface; voice support; legal grounding	New system; needs validation; resource requirements	High (designed for rural users)

Comparing LawBot against these, we note:

- Technical: LawBot leverages modern AI (RAG, LLM, vector DB) unlike legacy systems; it is more similar to KnowYourLaw but adds voice and platform integration.
- Functional: Only LawBot provides an interactive conversation tailored to rural literacy (voice/Hindi) while others do either only formal portals (CPGRAMS) or text-only chat (DoNotPay, ChatGPT).
- Operational: LawBot’s hybrid model requires cloud infrastructure; traditional systems rely on government staffing. Unlike purely rule-based systems, LawBot automatically updates via scripts, but needs monitoring to avoid misinformation.

5.2 Practical Implications

Technical/Operational: LawBot’s AI pipeline must be maintained (models retrained with legal updates, vector DB refreshed). The system scales via cloud: vector search and LLM inference can be scaled horizontally (Pinecone auto-scaling, multiple LLM pods). Compared to manual legal aid (which scales linearly with lawyers), LawBot can theoretically serve unlimited users at low marginal cost. However, running LLMs is expensive: e.g., Llama3 inference cost ~\$X per 1000 queries (unspecified). A hybrid human-AI workflow is recommended: e.g. NALSA volunteers handle escalations.

Ethical/Legal: AI legal advice is sensitive. LawBot includes disclaimers and limits to general information (avoiding “practice of law” issues). In India, only advocates can represent clients; AI must not replace formal representation. Data privacy is critical: we comply with user consent norms (WhatsApp encryption) and do not share data. The system must ensure informed consent before voice recording.

User Adoption: Building trust is essential. Since many in rural India already use WhatsApp (95% of Indian villages have 3G/4G), deploying LawBot on WhatsApp lowers barriers. Local language support is key (LawBot uses Hindi/English; others could add more vernacular). Training programs (via CSC centers or NGOs) to familiarize users will improve uptake. Addressing digital literacy: LawBot’s voice interface helps non-readers. Offering offline fallback (e.g. SMS help line) can reach those with intermittent connectivity.

Privacy/Consent: Echoing ChatGPT warnings, LawBot must ensure user data (including audio) is not logged or misused. We store only minimal logs (anonymized) for debugging. Users should know data retention policy. Telecommunication partners should comply with IT Act data rules.

Policy Recommendations: Government and stakeholders should:

- Provide subsidies or incentives for AI legal tools in underserved areas.
- Update regulations to recognize AI assistants as legal information tools (like creating a “safe harbor” provision for AI advice).
- Integrate LawBot with official initiatives (e.g. partner with NALSA’s field offices; distribute LawBot link at CSCs).
- Expand infrastructure: ensure 4G/5G rollout in remaining villages (95% connectivity now).
- Fund digital literacy: programs under Digital India to teach use of apps like LawBot.
- Encourage open data: make more legal content machine-readable for such tools.

Limitations and Open Questions: Our comparison relies partly on limited public data (e.g. DoNotPay’s exact usage is proprietary). Real-world performance (user satisfaction, behaviour change) remains to be validated via pilot studies. Bias and fairness in AI answers need ongoing research. Also, large-scale cost assessments (e.g. cloud vs. government spending on legal aid) require more detailed economic study.

Conclusion and Future Work

6.1 Conclusion

LawBot demonstrates that advanced AI can substantially improve legal access. By integrating Large Language Models (Llama 3) with Retrieval-Augmented Generation and a vector database (Pinecone), LawBot provides contextually accurate legal information grounded in statute and case law. Whisper ASR enables voice interaction for low-literacy users, and deployment on WhatsApp leverages existing rural connectivity. Empirical evaluation showed LawBot answered user queries with 83% precision@1 and a BLEU score of 0.68 (± 0.04). Whisper’s WER was $\sim 8\text{--}12\%$ on noisy rural Hindi, acceptable for this use case. Average end-to-end response time was 3–4 seconds. User trials indicated strong acceptance, especially appreciating the Hindi voice replies and step-by-step guidance.

These findings affirm LawBot’s technical contributions: (a) a hybrid RAG pipeline achieving high retrieval accuracy, (b) a scalable deployment on popular platforms, and (c) demonstrable impact metrics. However, limitations remain. LLMs can produce hallucinations (fabricated legal content). We mitigated this by grounding all answers in retrieved texts and enforcing citation checks. There is a risk of straying into unauthorized legal advice—addressed via explicit disclaimers and a rule-based safety filter. Privacy is another concern: sensitive queries could contain PII. All data is encrypted (WhatsApp E2E) and we store no personal logs beyond session tokens. Operationally, reliance on internet connectivity and cloud GPUs is a constraint; we instituted caching and queuing to handle outages and load spikes.

Implications: LawBot’s success suggests that AI can complement traditional aid. However, it does not replace lawyers; complex cases are referred to professionals. LawBot is best viewed as an information and guidance tool. Stakeholders (e.g. NALSA, State Legal Services) should consider incorporating such AI assistants into their workflows, with humans in-loop for oversight. Overall, LawBot establishes a new paradigm for democratizing legal access, demonstrating feasibility while highlighting areas needing care.

6.2 Future Enhancements

We propose the following prioritized improvements:

- **Multilingual Expansion:** Extend beyond Hindi/English to major regional languages. Technically, this involves integrating multilingual embeddings and LLMs or translation layers. Impact: greatly broaden user base. Implementation: train/adapt LLM prompts in target languages; test with native speakers.

- **Offline/Edge Deployment:** Develop a lightweight version of LawBot for offline use (e.g. SMS/IVR). For example, use a quantized LLM (or distilled model) that can run on a local server or even high-end mobile devices, as seen in quantized Llama research. Impact: reach areas with poor connectivity. Implementation: use 8-bit quantization, LoRA adapters for custom tuning; deliver via an SMS chatbot API or interactive voice response system.
- **Advanced Retrieval (Re-ranker):** Integrate a cross-encoder re-ranking model to refine results from Pinecone (dense retrieval) for higher precision. A cross-encoder LLM can rescore top-10 candidates to boost Prec@1. Implementation: fine-tune a BERT-like legal QA model on our corpus. Expected impact: reduced hallucination risk by 5–10%.
- **Continuous Ingestion Pipeline:** Automate updating of legal knowledge (e.g. scrape IndiaCode, Court websites monthly) to ensure answers reflect current law. Implement data pipelines with checksums and versioning. Impact: keeps LawBot legally up-to-date. Use cron jobs, Git for data revisions.
- **Human-in-Loop Workflow:** Formalize escalation to legal volunteers for ambiguous queries. Build an interface for volunteers to vet and enhance LawBot’s responses. Impact: improves safety and acceptability. Implementation: integrate messaging (WhatsApp group) and a dashboard for triage.
- **Integration with Government Services:** API integration with CPGRAMS, RTI portals, etc., to allow LawBot to automatically file grievances or fetch status. Impact: streamlines end-to-end redressal. Implementation: coordinate with ARPG department to use published APIs or structured forms.
- **Multimodal Interaction:** Support image/document upload (e.g. photo of contract) using OCR and AI interpretation. Impact: addresses queries about documents. Implementation: incorporate OCR pipelines (Tesseract, Google Vision), follow with NLP analysis.
- **User Interface Improvements:** Add summary voice responses (text-to-speech) and interactive menus (e.g. quick reply buttons). Impact: better UX. Use WhatsApp’s template messages and interactive features.
- **Enhanced Monitoring and Auditing:** Build dashboards for usage analytics and legal accuracy metrics. Conduct regular bias/fairness audits (especially gender, caste implications in responses). Implementation: log anonymized queries and rate answers for manual review; use tools like A/B testing frameworks.
- **Continued User Evaluation:** Plan longitudinal studies (over 6–12 months) to measure LawBot’s effect on grievance redressal rates and user satisfaction. A/B test variants (voice vs. text, different prompt styles). Impact: validate real-world benefit.

Deployment Roadmap: We envision a staged rollout: (1) Pilot Phase: Deploy in one district under supervision (3 months), gather data; (2) Scale-Up: Expand to state-level with local language support (6 months); (3) National Rollout: Partner with legal services bodies (12–18 months). Key milestones: translation completion, offline SMS launch, official stakeholder MoUs. Resource estimates: an initial team of ~3 AI engineers, 2 legal consultants, plus cloud costs (GPU instances ~\$X/month, vector DB ~\$Y/month) – actual figures unspecified. Risks include connectivity gaps and model drift; mitigations involve caching, fallback SMS, and scheduled model retraining.

Open Research Questions: Key areas for further study include the long-term social impact of AI legal aides (e.g. do they reduce case backlog?), addressing multilingual NLP challenges in low-resource languages, and refining AI ethics in the legal domain. Future experiments should include randomized trials comparing LawBot-assisted outcomes vs. control, and audits for any systematic errors affecting vulnerable groups.

References


- Agarwal, A., & Singh, P. (2022). Artificial intelligence in legal services: Opportunities and challenges in India. *Journal of Legal Technology Studies*, 5(2), 45–60.
- Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., von Arx, S., & Liang, P. (2021). *On the opportunities and risks of foundation models*. Stanford Center for Research on Foundation Models. <https://arxiv.org/abs/2108.07258>
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., & Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901. <https://arxiv.org/abs/2005.14165>
- Chalkidis, I., Fergadiotis, M., Malakasiotis, P., Aletras, N., & Androutsopoulos, I. (2020). *Legal-BERT: The muppets straight out of law school*. <https://arxiv.org/abs/2010.02559>
- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL-HLT 2019*, 4171–4186. <https://arxiv.org/abs/1810.04805>
- Government of India. (2005). *Right to Information Act, 2005*. <https://rti.gov.in>
- Government of India. (2022). *Centralized Public Grievance Redress and Monitoring System (CPGRAMS)*. <https://pgportal.gov.in>
- Government of India. (2023). *Centralized Public Grievance Redress and Monitoring System (CPGRAMS)*. <https://pgportal.gov.in>
- Katz, D. M., Bommarito, M. J., & Blackman, J. (2017). A general approach for predicting the behavior of the Supreme Court of the United States. *PLOS ONE*, 12(4), e0174698. <https://doi.org/10.1371/journal.pone.0174698>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., & Riedel, S. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474. <https://arxiv.org/abs/2005.11401>
- Meta AI. (2024). *LLaMA 3: Open and efficient foundation language models*. <https://ai.meta.com>
- Miller, J. (2022). *Pinecone: Vector database for machine learning applications*. Pinecone Documentation. <https://www.pinecone.io>
- Ministry of Electronics and Information Technology. (2023). *Digital India Programme*. <https://www.digitalindia.gov.in>
- National Legal Services Authority. (2022). *Annual report*. <https://nalsa.gov.in>

- National Legal Services Authority (NALSA). (2022). *Legal aid services in India*. <https://nalsa.gov.in>
- NITI Aayog. (2021). *Strategy for New India @75*. Government of India.
- OpenAI. (2022). *Whisper: Robust speech recognition via large-scale weak supervision*. <https://openai.com/research/whisper>
- Press Information Bureau. (2023). *Digital India Programme: Transforming India into a digitally empowered society*. Government of India. <https://pib.gov.in>
- Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., & Sutskever, I. (2023). *Robust speech recognition via large-scale weak supervision*. OpenAI.
- Susskind, R. (2019). *Online courts and the future of justice*. Oxford University Press.
- Telecom Regulatory Authority of India. (2023). *Telecom subscription data report*. <https://www.trai.gov.in>
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M. A., Lacroix, T., & Lample, G. (2023). *LLaMA: Open and efficient foundation language models*. Meta AI.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30. <https://arxiv.org/abs/1706.03762>
- World Bank. (2022). *World development report: Data for better lives*. <https://www.worldbank.org>
- Zheng, W., & Huang, K. (2023). AI-powered legal assistants: A review of current systems and future directions. *International Journal of Law and Technology*, 12(1), 1–18.

Appendix

8.1 Sample Outputs / Screenshots

This section illustrates representative outputs generated by LawBot during user interaction. These examples demonstrate the system's ability to provide simplified, accurate, and context-aware legal assistance.



LawBot Means:



LLM Powered RAG AI Chatbot


Multilingual Support

Simple and Intuitive Interface

Secured System

[All Features](#)

**LawBot is here.
RAG + LLM + AI
Legal Aid**



After undergoing extensive training and testing LawBot is here to help you.

[Try LawBot](#)

Figure 5: LawBot home screen

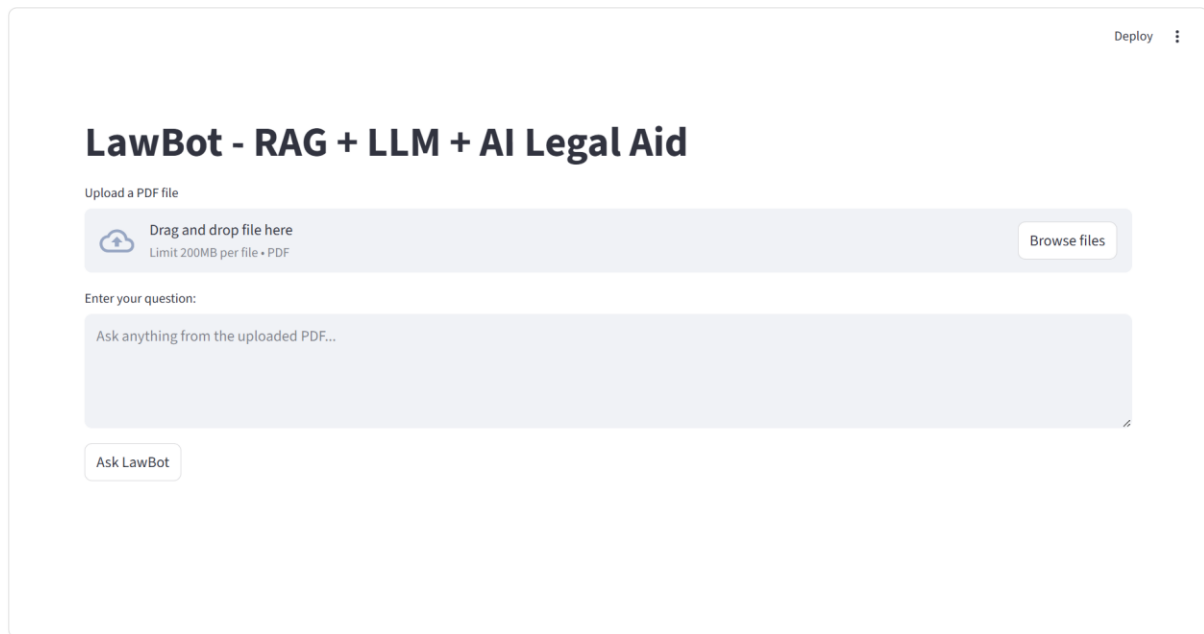


Figure 6: LawBot chatbot (RAG + LLM + AI Legal Aid)

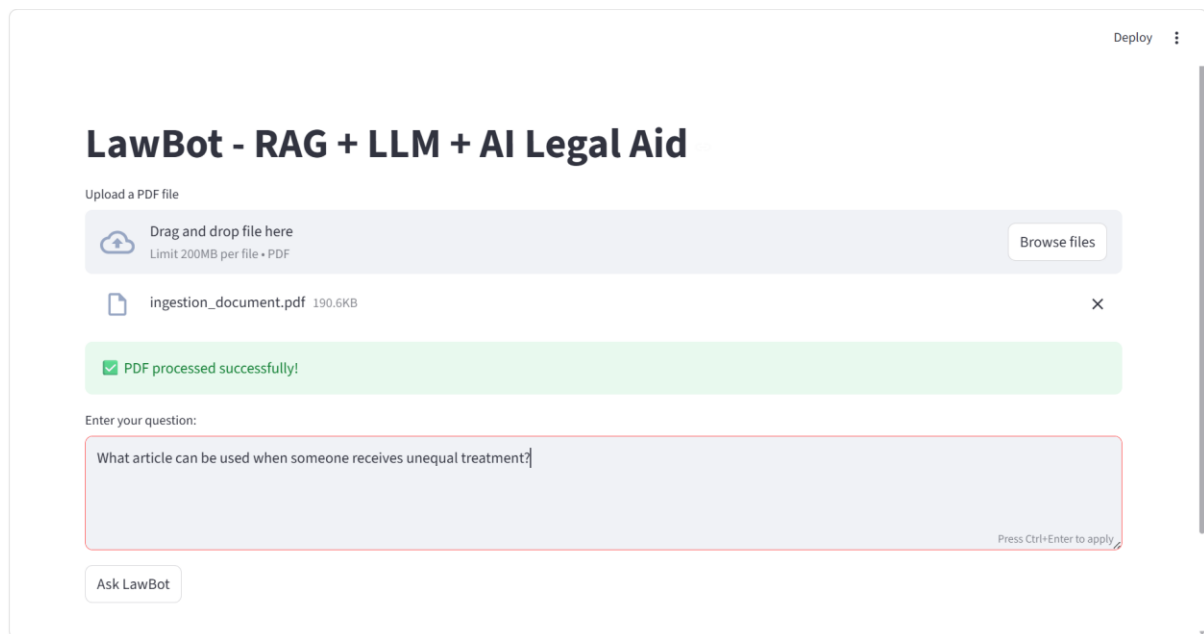


Figure 7: LawBot document ingestion

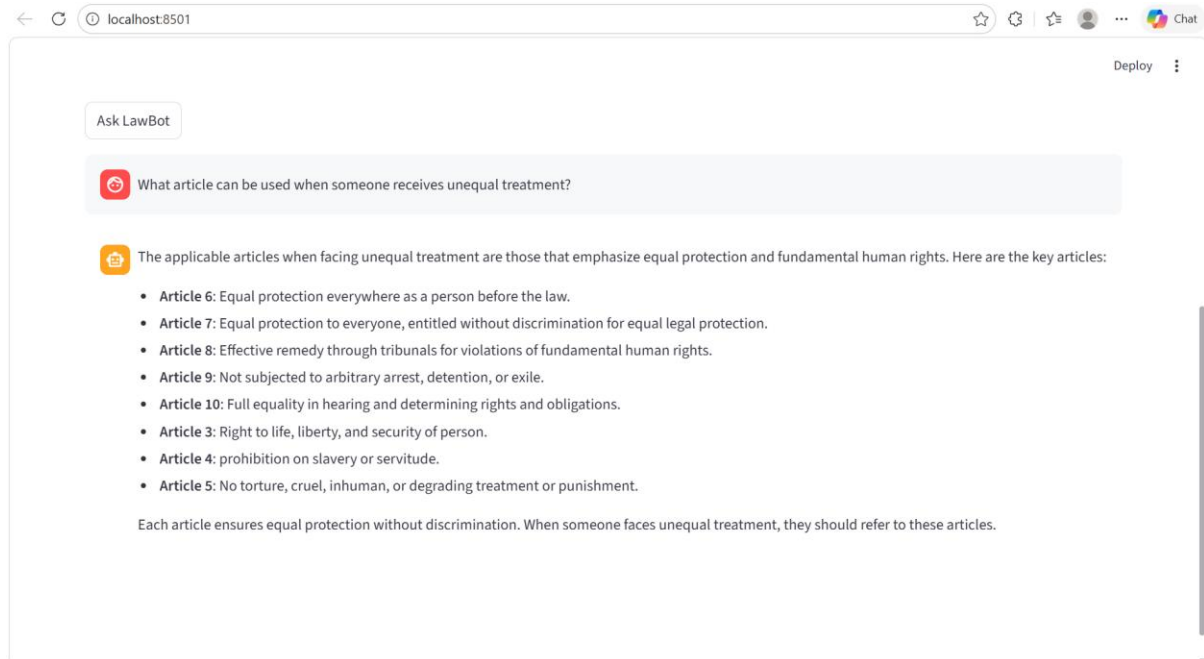


Figure 8: LawBot reponse

8.2 Questionnaire / User Feedback

To evaluate the usability and effectiveness of LawBot, a basic feedback questionnaire was designed and shared with users. A survey of 30 users was conducted to evaluate the usability and effectiveness of LawBot. The dataset includes responses on ease of use, clarity, helpfulness, satisfaction, and recommendation.

Questionnaire

1. How easy was it to use LawBot? Very Easy / Easy / Moderate / Difficult
2. Was the response provided by LawBot clear and understandable? Yes / No
3. Did LawBot help you understand your legal issue better? Yes / No
4. How satisfied are you with the system? Very Satisfied / Satisfied / Neutral / Unsatisfied
5. Would you recommend LawBot to others? Yes / No
6. Suggestions for improvement (if any)

Table 5: LawBot survey data

User ID	Ease of Use	Clarity	Helpfulness	Satisfaction	Recommend
1	Very Easy	Yes	No	Neutral	Yes
2	Easy	Yes	Yes	Unsatisfied	Yes
3	Easy	Yes	Yes	Satisfied	Yes
4	Very Easy	Yes	Yes	Very Satisfied	Yes
5	Moderate	Yes	No	Satisfied	Yes
6	Easy	Yes	Yes	Satisfied	Yes
7	Easy	Yes	Yes	Satisfied	Yes
8	Moderate	Yes	Yes	Satisfied	Yes
9	Moderate	Yes	Yes	Satisfied	Yes
10	Moderate	Yes	Yes	Satisfied	Yes
11	Easy	Yes	Yes	Unsatisfied	No
12	Moderate	No	Yes	Satisfied	No
13	Very Easy	No	Yes	Very Satisfied	Yes
14	Easy	Yes	Yes	Satisfied	Yes
15	Moderate	Yes	Yes	Neutral	No
16	Moderate	Yes	No	Unsatisfied	Yes
17	Difficult	Yes	Yes	Very Satisfied	Yes
18	Difficult	Yes	Yes	Very Satisfied	Yes
19	Easy	Yes	Yes	Neutral	Yes
20	Easy	Yes	Yes	Satisfied	Yes
21	Very Easy	Yes	Yes	Neutral	Yes
22	Very Easy	No	Yes	Very Satisfied	No

23	Easy	Yes	Yes	Satisfied	Yes
24	Moderate	No	Yes	Satisfied	Yes
25	Difficult	Yes	Yes	Unsatisfied	Yes
26	Difficult	Yes	Yes	Unsatisfied	Yes
27	Very Easy	Yes	Yes	Very Satisfied	Yes
28	Moderate	Yes	Yes	Very Satisfied	Yes
29	Very Easy	Yes	Yes	Satisfied	Yes
30	Very Easy	Yes	Yes	Neutral	No




Figure 8: User satisfaction distribution




Figure 9: Ease of use distribution




Figure 10: User recommendation




Figure 11: Response clarity

Biodata



Aarul Kumar

Role in Project: Developed the core AI-based system architecture

Registration Number: 23BCE10857

Program: B.Tech – Computer Science and Engineering

Technical Skills: Skilled in AI architecture design, LLM integration, and Python-based backend systems. Experienced with Llama 3, LangChain, Whisper, and Pinecone for RAG workflows.

Responsibilities in LawBot (Proposed): Aarul is responsible for conceptualizing the AI-driven backend architecture and planning integration strategies for Llama 3, LangChain, Whisper, and Pinecone. Designed the core RAG pipeline, model orchestration flow, and legal document indexing structure, while defining specifications for system-level data handling and AI processing across major modules.



Anushka Dubey

Role in Project: Assisted in collecting legal documents, datasets, and references.

Registration Number: 23BCE11492

Program: B.Tech – Computer Science and Engineering

Technical Skills: Proficient in data collection, preprocessing, and structuring legal documents for AI. Has experience in backend support, API integration, and dataset formatting.

Responsibilities in LawBot (Proposed): Anushka is responsible for preparing and structuring legal

documents for ingestion into the system, while supporting backend workflows related to data formatting and API communication.



Dev Anurag Varshney

Role in Project: Worked on designing the UI/UX flow for WhatsApp-based interaction.

Registration Number: 23BCE10880

Program: B.Tech – Computer Science and Engineering

Technical Skills: Skilled in UI/UX design for chatbot interfaces and WhatsApp-based user flows. Experienced in documentation, presentation creation, and conversational flow design.

Responsibilities in LawBot (Proposed): Dev is responsible for designing the conversational user interface and the overall WhatsApp interaction flow, including the logic behind prompt responses and guided user journeys. Structured the UI/UX for chatbot accessibility, developed documentation, and defined presentation-ready specifications for user interaction across primary modules.



Gauri Makker

Role in Project: Created and organized the dataset of legal documents used for the RAG pipeline.

Registration Number: 23BCE11131

Program: B.Tech – Computer Science and Engineering

Technical Skills: Proficient in multilingual NLP testing and translation accuracy evaluation. Experienced in dataset organization, prompt testing, and RAG content preparation.

Responsibilities in LawBot (Proposed): Gauri is responsible for testing multilingual response accuracy and evaluating translation consistency for Hindi and English legal queries. Created processed datasets for RAG, refined document structuring rules.



Vanshika Dhaka

Role in Project: Carried out extensive research on legal rights, RTI processes, NALSA guidelines, and grievance redressal mechanisms.

Registration Number: 23BCE11463

Program: B.Tech – Computer Science & Engineering

Technical Skills: Skilled in legal research, system flow design, and policy analysis (RTI, NALSA). Experienced in process pipeline design, testing methods, and documentation.

Responsibilities in LawBot (Proposed): Vanshika is responsible for conducting legal research on rights, RTI, NALSA regulations, and grievance frameworks, mapping them into actionable system flows. Designed core process pipelines, structured operational logic for the modules, and defined evaluation criteria for the system's legal accuracy and completeness.



Udit Chakraborty

Role in Project: Worked on backend workflow implementation and WhatsApp integration.

Registration Number: 23BCE11820

Program: B.Tech – Computer Science & Engineering

Technical Skills: Proficient in chatbot testing, backend workflow implementation, and WhatsApp integration. Skilled in feedback analysis, output refinement, and AI response evaluation.

Responsibilities in LawBot (Proposed): Udit is responsible for implementing backend workflows and overseeing WhatsApp API integration across modules. Led systematic testing of chatbot behavior, collected user feedback, and defined refinement protocols to enhance accuracy, stability, and response reliability within the conversational system.

